

BatSignal: System Design Document

Joe Moraal
moraalj@wit.edu

Zach Thornton
thorntonz@wit.edu

Bryan Young
youngb2@wit.edu

Computer Science 2015
Wentworth Institute of Technology

June 28, 2015

Contents

1	Introduction	3
1.1	Purpose and Scope	3
1.2	Project Executive Summary	3
1.2.1	System Overview	3
1.2.2	Design Constraints	4
1.2.3	Future Contingencies	4
1.3	Points of Contact	4
1.4	Project References	5
1.5	Glossary	6
1.5.1	System Specific Definitions	6
1.5.2	Technical Definitions	6
1.5.3	Industry Definitions	8
1.6	Document Organization	8
2	System Architecture	9
2.1	System Hardware Architecture	9
2.1.1	Control Node	9
2.1.2	Sensor Node	9
2.2	System Software Architecture	9
2.3	Internal Communications Architecture	10
2.3.1	Wireless Mesh Network	10
3	Human-Machine Interface	10
3.1	Inputs	10
3.2	Outputs	11
4	Detailed Design	11
4.1	Hardware Detailed Design	11
4.1.1	Raspberry Pi 2	11
4.1.2	Wi-Pi WLAN Module	12
4.1.3	Microphone	12
4.2	Software Detailed Design	12

1 Introduction

1.1 Purpose and Scope

This document describes the hardware and software components of the BatSignal distributed sensor network. This document is intended for use by developers implementing BatSignal.

1.2 Project Executive Summary

The system is designed as a rapid response alert system capable of identifying emergencies and reporting their location. The system passively captures audio from the sensors and analyzes it for keywords or phrases. When the system detects a match it dispatches an email to a list of administrators and displays a notification on the administration console.

The system is designed to be scalable according to the needs of the location of installation. Controller nodes are installed at or near administrative areas with sensor nodes installed in patient rooms, inhabited spaces, common areas, etc. Messages propagate through the BatSignal mesh network allowing nodes to communicate with the controller node despite the often significant distance between them.

1.2.1 System Overview

The system is divided into two types of nodes. These nodes form a mesh network that relays data from the first type of node, called sensor nodes, to the second type of node, called the controller node.

Python modules installed on the sensor nodes passively read input from a microphone. The input from the microphone is fed into Python's speech recognition module. The audio input is then sent to the controller node over the mesh network as JSON.

Python modules installed on the controller node passively receive the sensor's data. It then parses through the text looking for alert phrases. Upon recognizing an alert phrase an email is sent to a list of administrators containing the plain-text of the audio capture and the identifier for the sensor node that captured it.




Figure 1: BatSignal system overview

1.2.2 Design Constraints

The BatSignal distributed sensor network will consist of numerous Raspberry Pi boards distributed across a physical area. This network has many constraints that must be considered. These constraints are as follows:

- **Maximum Node Separation Distance** - The nodes must be close enough to communicate effectively with one another in order to form a reliable mesh network. This maximum effective communication distance relies on the capabilities of the Wi-Fi WLAN module and the B.A.T.M.A.N. routing protocol. This distance has not been determined yet.

The communication range is also affected by physical objects in the installation environment, which may reduce the base measurement during implementation.

- **Sensor per Controller Ratio** - Due to the physical limitations of the controller node's processing power and each Wi-Fi WLAN module's transmission rate, coupled with the mesh network's maximal flow limits, there is a finite number of sensors that should be paired with any single controller node. This ratio has not been determined yet.

1.2.3 Future Contingencies

The BatSignal distributed sensor network relies heavily on the ability for nodes to be physically separate from one another, and still networked to the Internet. This is achieved by leveraging the B.A.T.M.A.N. routing protocol. Should the B.A.T.M.A.N. protocol be inadequate for wireless communication however, alternatives such as project Meshnet may prove more usable.

BatSignal also relies heavily on the ability to convert audio captures to a textual representation. The project leverages Google speech-to-text APIs to perform text conversion. The reliability and feasibility of these APIs may require better alternatives. No such alternatives have yet been identified as contingencies.

1.3 Points of Contact

- **Physical Layout**
This role involves determining the optimal distance between nodes in the network and their placement within the environment. It is closely connected to the role for Installation but also requires benchmarking to find the limits and recommended distances for use in the documentation of the system.
 - primary: Zach Thornton
- **BatSignal Sensor Module**
This role involves designing and implementing the python module for the sensor nodes. This module will have the design detailed in section 4.2.
 - primary: Joseph Moraal
 - secondary: Bryan Young
- **BatSignal Control Module**
This role involves designing and implementing the python module for the controller node. This module will have the design detailed in section 4.2.
 - primary: Joseph Moraal
 - secondary: Bryan Young

- **Administration Console**
This role involves designing and implementing the administration console detailed in section 4.2.
 - primary: Zach Thornton
- **Documentation**
This role involves writing, proof reading, and compiling the documentation for the system. This documentation involves all documents asked for including this document as well as installation guides, performance analytics, etc.
 - primary: Bryan Young
 - secondaries: Joseph Moraal, Zach Thornton
- **Network Setup**
 - primary: Bryan Young
 - secondary: Zach Thornton
- **Performance Analysis**
This role involves benchmarking the limits of the system and recording them. These benchmarks include latency, throughput, and accuracy.
 - primary: Joseph Moraal
- **Installation**
This role involves the installation of the system within the test environments and the documenting of the steps involved from physical hardware setup to software setup and initial configuration.
 - primary: Zach Thornton
 - secondary: Bryan Young
- **Maintenance**
This role involves determining the steps required to perform normal maintenance of the system and documenting them. Examples of this are: documenting the administration console, and the controller node configuration. Documentation for the administration console includes: how to connect to the console and navigate to different portions within the console. Documentation on the controller node configuration includes how to ssh into the controller node, where to find the configuration files used by the system, and how to restart the system.
 - primary: Joseph Moraal
 - secondary: Zach Thornton

1.4 Project References

- **Raspberry Pi:** <https://www.raspberrypi.org/>
- **B.A.T.M.A.N. Routing Protocol:** <http://www.open-mesh.org/projects/batman-adv>
- **Python Speech Recognition Library:** <https://pypi.python.org/pypi/SpeechRecognition/>
- **Python JSON Encoder and Decoder:** <https://docs.python.org/3.3/library/json.html>
- **JavaScript Object Notation:** <http://json.org/>

1.5 Glossary

The glossary provides expansions for acronyms and abbreviations which appear within this document. It also provides definitions for terminology used within the document.

1.5.1 System Specific Definitions

System Specific Definitions	
Regular User	Any patient or person the system is designed to monitor
Administrative User	Any professional staff who are designated to receive email alerts from the system

1.5.2 Technical Definitions

Technical Definitions	
ad-hoc network	A decentralized type of network that does not rely on pre existing infrastructure such as routers
Amp	Amperage: A unit of measurement describing the number of electrons moving past a fixed point in a conductor per second
API	Application Programming Interface: A set of routines, protocols, and tools for building software
CPU	Central Processing Unit: The electronic component of a computer which executes instructions of a program by performing basic arithmetic, logical, control, and input/output operations
dB	Decibels: A logarithmic unit of measurement which expresses the ratio between two values of a physical quantity, herein referring to loudness
dBm	Decibel mili-watts: An abbreviation for the power ratio in decibels of the measured power reference to one mili-Watt.
DC	Direct Current: The unidirectional flow of electric charge
g	grams: The basic SI unit of measurement for mass
GHz	GigaHertz: A unit of measurement representing one billion Hertz
GPIO	General Purpose Input Output: The generic hardware input/output pin layout of an integrated circuit
GPU	Graphics Processing Unit: A multi-core computer chip designed to quickly alter and manipulate memory to accelerate graphical computation
Hertz:	A measurement of electromagnetic frequency, or the number of oscillation of the perpendicular electric and magnetic fields per second
I/O	Input/Output: The communication between two information processing systems

Technical Definitions	
kHz	kiloHertz: A unit measurement representing one thousand Hertz
mA	MiliAmps: A unit of measurement representing one one-thousandth of an Amp
MBps	MegaBits per second: A unit of measurement representing the data transfer speed as one million bits per second
mesh network	A network topology. See complete description in section 2.3.1
MHz	Mega-Hertz: A unit of measurement representing one million Hertz
MiB	MebiByte: An alternative measurement for what is pejoratively referred to as a MegaByte, meaning 2^{20} bytes
MicroSD	A miniaturized version of the SD card
mm	Mili meter: A unit of measurement representing one one-thousandth of a meter
SD card	Secure Digital Card: A physical memory device capable of digital information storage
SDRAM	Synchronous Dynamic Random Access Memory: Dynamic random access memory synchronize with the system bus
SoC	System on a Chip: An integrated circuit that integrates all components of a computer or other system into a single chip
V	Volts: The SI unit for electromotive force, the difference of potential that would drive one ampere of current against one ohm resistance
W	Watts: A derived SI unit of power expressing the rate of energy conversion or transfer with respect to time
WLAN	Wireless Local Area Network: A wireless computer network that links two or more devices using a wireless distribution method within a limited area

1.5.3 Industry Definitions

Industry Definitions	
ARM	A family of instruction set architectures
B.A.T.M.A.N	Better Approach to Mobile Ad-hoc Networking
Flask	A Python library used for backend Web Development
Google	A company specializing in Internet Technologies
IEEE	Institute of Electrical and Electronics Engineers
JSON	JavaScript Object Notation, a human readable format used to transmit data objects.
Project Meshnet	An alternative ad-hoc network protocol
PyAudio	A Python library used for cross platform audio I/O
Python	High level programming language
Raspberry Pi	A budget single board computer
SI	The international system of units which defines basic units of measurements
SpeechRecognition	A Python library used for accessing the Google Speech Recognition API
USB	Universal Serial Bus: Industry standard defining cables, connectors, and communication protocols used in a bus for connection, communication, and power supply between electronic devices
WEP	Wired Equivalent Privacy, a network security protocol
Wi-Pi	A wireless adapter for the Raspberry Pi
WPA	Wi-Fi Protected Access, a network security protocol
WPA-PSK	Wi-Fi Protected Access Pre-Shared Key, a network security protocol

1.6 Document Organization

In the following sections this document will define the overall system architecture followed by more detailed hardware, software, and communication architectures. These sections will be followed by the specifications for the system interface, both input and output. The final section will then go into explicit detail about each hardware and software component present within the system. Finally, the document ends with appendices containing reference or additional material.

2 System Architecture

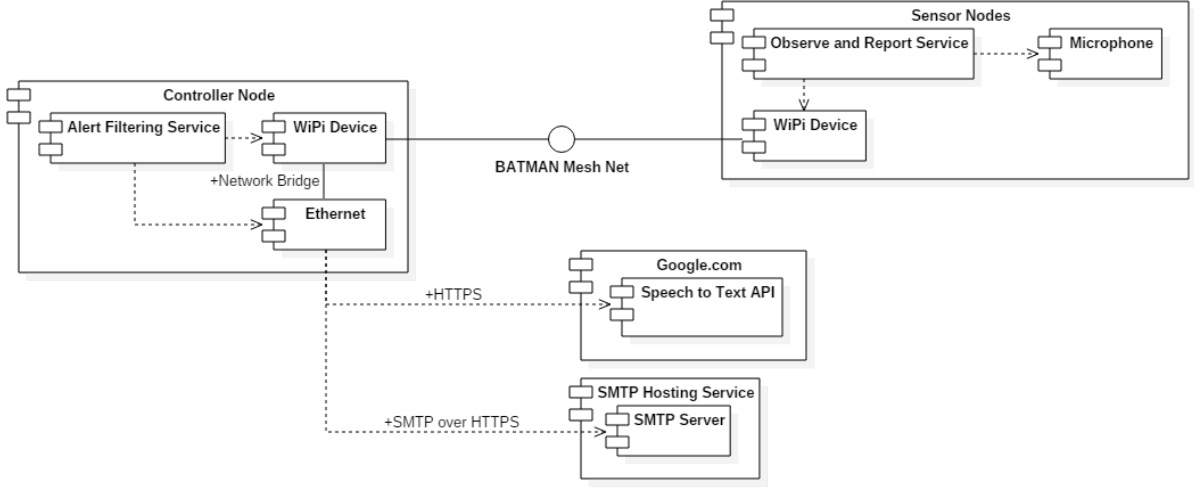


Figure 2: BatSignal system architecture

2.1 System Hardware Architecture

The BatSignal distributed sensor network consists of two types of hardware nodes, consisting primarily of a Raspberry Pi model 2 board. These boards are inexpensive and small computers which will enable the BatSignal network to perform processing and decision making tasks based on input to the system via connected sensors.

2.1.1 Control Node

The BatSignal controller node consists of the Raspberry Pi model 2 board outfitted with a Wi-Pi WLAN module which enables wireless networking capabilities. These nodes also utilize the on board Ethernet adapter to enable a wired network connection to the Internet.

2.1.2 Sensor Node

The BatSignal sensor node consists of the Raspberry Pi model 2 board outfitted with a Wi-Pi WLAN module which enables wireless networking capabilities. The sensor nodes are also outfitted with a USB microphone which enables the passive audio recording used to monitor the environment.

2.2 System Software Architecture

The system has two types of nodes in the mesh network. The first type of node is a sensor, and the second is the controller node. There is only one controller node per complete network. The sensor nodes and controller node run different python modules which dictate their behavior.

Sensor nodes passively wait for input from the attached microphone. Using Python's speech recognition library the audio input is converted into text. It is then sent across the mesh network to the controller node.

The controller node passively waits for input from the sensor nodes. Upon receiving a message it searches the text for alert phrases that are defined in a configuration file and read upon startup. When an alert phrase is identified, an email is composed with the full text of the

message and sent to the administrators. The administrators are defined by email address in a configuration file and read upon startup. The controller node is the only node connected directly to the Internet.

2.3 Internal Communications Architecture

The BatSignal distributed sensor network must be capable of communication between nodes, which will be carried out over a wireless mesh network.

2.3.1 Wireless Mesh Network

A mesh network is a topology in which each node connected to the network acts as a relay for data being passed through the network. Data passing through the network may be propagated using either a flooding or a routing technique. Not every node within the network must be connected to every other node on the network. The ability to propagate data by relaying it from node to node allows the network to be both distributed and fault tolerant. The simple fact that the mesh network is being implemented using wireless network adapters makes the mesh network a wireless mesh network.




Figure 3: A fully connected, and a partially connected mesh network.

3 Human-Machine Interface

The BatSignal distributed sensor network is designed to work transparently to the user. The system is designed to passively record audio from the surrounding environment. It does not require any interaction from the user to activate and maintenance can be performed remotely. Input from administrators to maintain the system is performed through remote access to the controller nodes via ssh and through the administration console. The system outputs email messages upon detecting an alert phrase. Email messages are sent to the list of administrators and contain the full text of the triggering input with the identifier of the node that originally detected the alert phrase.

3.1 Inputs

Input to the system comes in three forms; configuration files, the administration console, and ambient verbal audio.

Configuration files are read by each node on startup. These files are stored on the controller node

in a subdirectory of the installation folder named "config". The files are stored as "phrases.json" and "admins.json". They each contain a single json object with a list of strings. The list of string for "phrases.json" are the help phrases for the system, each of which is a single English word. The list of strings for "admins.json" are email addresses for the administrators.

The input to the sensor nodes is through verbal audio signals recorded using the attached microphone. The controller node receives input in the form of text messages relayed across the mesh network from the sensor nodes.

The administration console may be used to modify system configuration.

3.2 Outputs

There are three forms of output in the system: internal output from the sensor nodes, email messages to the administrators, and status information available through the administration console.

The sensor nodes produce text messages containing the converted audio input recorded. These messages are relayed over the wireless mesh network to the controller node for processing. The controller node dispatches an email to the administrators containing the plain-text captured input and the node identifier of the sensor which captured the alert.

The administration console displays information concerning the current state of the system. The state includes sensor node information such as node uptime, node status, and hostnames. The console will also display any alerts that were emailed.

4 Detailed Design

4.1 Hardware Detailed Design

4.1.1 Raspberry Pi 2

Both versions of BatSignal nodes target the Raspberry Pi model 2 board. These systems have the following capabilities:

Raspberry Pi 2 Specifications	
Cost:	\$35 USD
SoC:	Broadcom BCM2836
CPU:	900MHz quad-core ARM Cortex-A7
GPU:	Broadcom VideoCore IV, OpenGL ES 2.0, OpenVG 1080p30 H.264 high-profile encode/decode
Memory (SDRAM)iB:	1024 MiB
USB 2.0 Ports:	4 (via intergrated USB hub and LAN9512)
Onboard Storage:	Micro Secure Digital / MicroSD slot
Onboard Network:	10/100 wired Ethernet RJ45
Real-time Clock:	None
Power Ratings:	650 mA, (3.0 W)
Power Source:	5 V (DC) via Micro USB type B or GPIO header
Size:	85.0mm x 56.0 mm x 17mm
Weight:	40g

4.1.2 Wi-Pi WLAN Module

Wi-Pi WLAN Module Specifications	
Cost:	\$15.52
Physical Interface:	USB 2.0
Wireless Standards:	IEEE 802.11n Backward compatible with IEEE 802.11g and IEEE 802.11b
Transmission Speed:	11b: 1/2/5.5/11 Mbps 11g: 6/9/12/18/24/36/48/54 Mbps 11n: up to 150 Mbps
Frequency Range:	2.4 to 2.4835 GHz
Working Channel:	1 to 13
Transmit Power:	20dBm (max)
Security Features:	WPA-PSK/WPA2-PSK WPA/WPA2 64/128/152 bit WEP Encryption

4.1.3 Microphone

Microphone Specifications	
Frequency Response:	50Hz - 18kHz
Polar Pattern:	Directional
Resolution:	16 Bit/44.1 kHz
Sensitivity:	56dB @ 1kHz
Output Voltage:	1.20V
Input:	USB

4.2 Software Detailed Design

There are three separate software components to the BatSignal distributed sensor network. The sensor nodes, controller node, and administration console.

The sensor nodes have three python modules installed on them. Two of the modules are from the official python libraries: PyAudio, and SpeechRecognition. These modules are used by the third module: the BatSignal sensor module. The BatSignal sensor module passively listens for input from the microphone. A callback is used to respond when the PyAudio module detects input from the microphone. The callback converts the audio input text using the SpeechRecognition module. The converted text is then logged and sent across the wireless mesh network to the controller node. The sensor network also adjusts for ambient noise. The system adjusts for ambient noise periodically using the PyAudio module. The time between adjustments is configurable through the administration console.

The controller node module, upon startup, reads from a series of configuration files stored in a subdirectory of it's installation folder. These configuration files are "phrases.json" and "admins.json", and both are simple JSON files containing a single JSON object and a list of strings. Using python's JSON decoder these files are parsed and stored by the module. The BatSignal control module then switches into a passive listening mode from which it waits for messages from the sensor nodes. Upon receiving a message from a sensor node it parses the JSON of the message into its components. The JSON message contains the identifier of the dispatching sensor node, the time the message was sent, and the contents of the message. The contents of the message are parsed using a simple string comparison method. The string is split using the delimiters of whitespace and control characters (commas, periods, etc.). Each word in

the split string is then compared to every alert phrase from "prases.json" Upon finding a match an email is composed containing the name of the sensor node the message originated from, the time the audio was recorded, and the full contents of the audio converted to text. The email is then mailed to all addresses listed in "admins.json".

The administration console will be the central configuration point for the BatSignal network. The backend for the console will be written using the python library Flask. The administration console will allow for the modification of the "phrases.json" and "admins.json" configuration files and will contain information related to the current BatSignal deployment. It will have the capability of performing diagnostic tasks as well as displaying current deployment information. These diagnostic tasks include restarting nodes, changing hostnames, and testing connectivity to all nodes. Deployment information that will be seen within the console includes hostnames, assigned IP addresses, and sensor status. The console will contain a panel that will detail the same alerts that have been emailed, and will be updated in real-time. Any nonfunctioning sensors will be visible within this panel.